

ASL Alphabet Image Recognition – Code Demonstration

Project Summary

This project classifies American Sign Language (ASL) letters (A–Z plus `del`, `space`, and `nothing`) using image data. A pre-trained **VGG16 model** is used for feature extraction, and custom dense layers are added to specialize the network for ASL classification.

Step-by-Step Code Explanation

1. Setup & Environment Configuration

- **Kaggle Authentication:** Copy the `kaggle.json` API key to `~/.kaggle/`.
 - **Download Dataset:** Use `kaggle datasets download -d grassknotted/asl-alphabet`.
 - **Unzip** the contents into a directory called `asl-alphabet`.
-

2. Dependencies

- TensorFlow/Keras (model building)
- OpenCV & PIL (image loading/processing)
- NumPy, Pandas (data handling)
- Plotly, Matplotlib (visualization)
- `imutils`, `sklearn`, `tqdm`, `glob`

3. Configuration & Seeding

- Hyperparameters are stored in a class `CFG`.
- `batch_size = 64`
- `img_height = 64`
- `img_width = 64`
- `epochs = 10`
- `num_classes = 29`

4. Label Creation & Metadata Building

- Uses `string.ascii_uppercase + ['del', 'nothing', 'space']` to define 29 class labels.
- Each label's directory is scanned for image paths and stored in a metadata DataFrame.

5. Dataset Splitting

- Uses `train_test_split` (stratified) to split metadata into:
 - 70% Training
 - 15% Validation
 - 15% Testing

6. Image Data Generators

- Uses `ImageDataGenerator(rescale=1./255)` for normalization.
 - Image generators load and batch the image files from disk efficiently.
-

7. Model Creation – VGG16 Transfer Learning

- **Base Model:** VGG16 (excluding top classifier).
 - **Custom Top Layers:**
 - `Flatten`
 - `Dense(256, relu) → Dropout`
 - `Dense(512, relu) → Dropout`
 - `Dense(29, softmax)`
 - The base VGG16 layers are frozen to retain learned features.
-

8. Model Compilation & Training

- **Loss Function:** Categorical Crossentropy
 - **Optimizer:** Adam
 - **Metrics:** Accuracy
 - Uses `ModelCheckpoint` to save the best model based on validation accuracy.
-

9. Evaluation

- Loads the saved model weights from disk.
 - Uses the test data to generate predictions.
 - Evaluates performance using a **confusion matrix**.
-

10. Feature Extraction & Visualization

- A secondary model is created to extract **512-dimensional features** from the last dense layer.
 - Uses **t-SNE** to reduce feature vectors to 2D.
 - Plots 2D feature space with labels to visualize class separability.
-

Demonstration Plan

Use this plan as a structure for your live or recorded demo:

1. **Introduction:** Briefly explain ASL and the goal of the project.
2. **Environment Setup:** Walk through mounting Google Drive, loading Kaggle API key, and installing dependencies.
3. **Preprocessing:** Show how metadata is created from directory structure.
4. **Model Building:**
 - Load VGG16
 - Add dense layers
 - Display `model.summary()`
5. **Training:** Explain loss, optimizer, and show the live training output.
6. **Evaluation:** Show the final test accuracy or confusion matrix.
7. **Visualization:** Display the t-SNE plot to explain feature embedding.